

Heap Sort

Refer to Chap 6 CLRS

Sakeena | 25-01-2022

Heap Sort


- Comparison based sorting algorithm.
 - Individual keys of the input are compared to one another to perform sorting.
- Time complexity \rightarrow like merge sort [$O(n \log n)$]
- In-place \rightarrow like insertion sort.
- Makes use of heap data structure.



Combines the better attributes of the two.

Recall

Binary Heaps




Heap DS - Recall

- Binary heaps can be represented using arrays. Consider an array A representing a max heap.
- The root of the heap is $A[1]$.
- Given index of any node, we can compute the index of its parent, left child or right child in the array A by using the following:

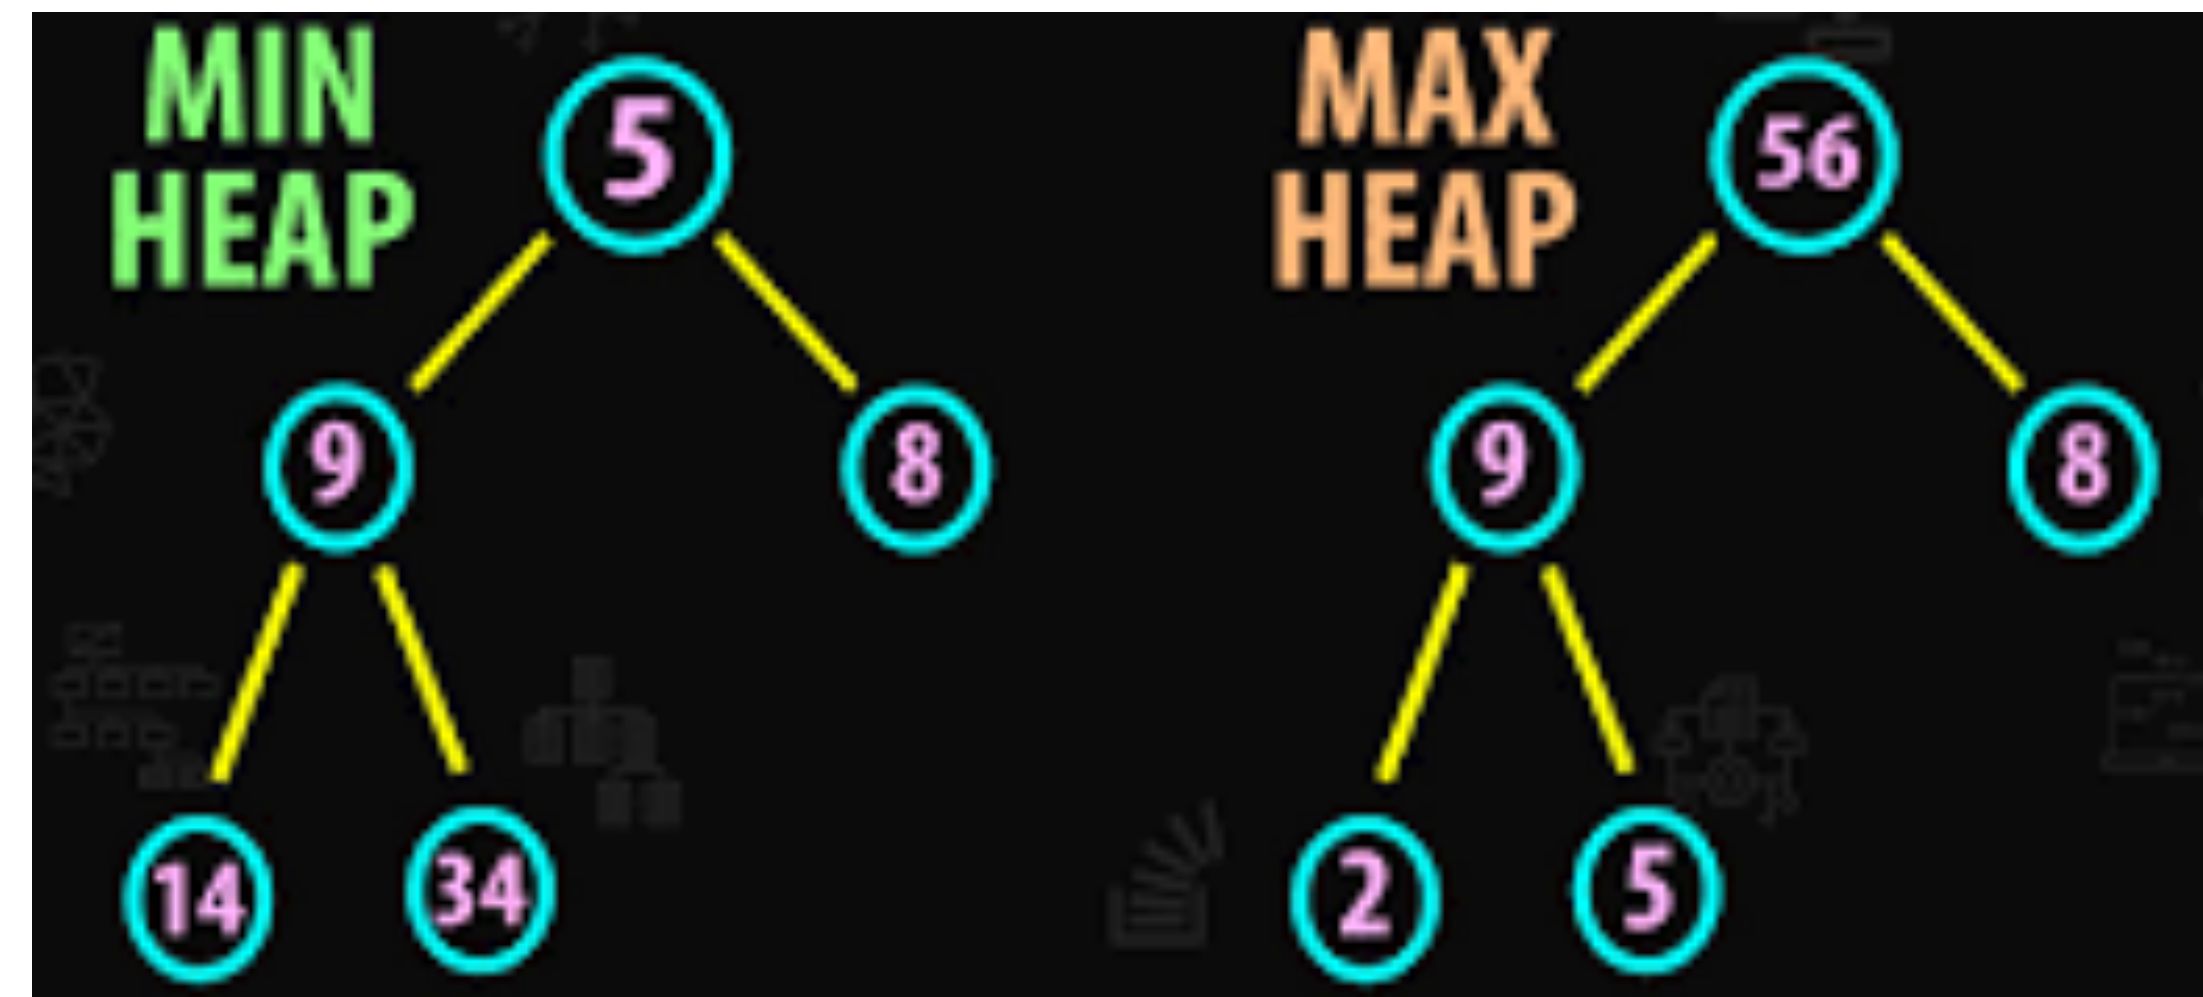
- $\text{Parent}(i) = \lfloor i/2 \rfloor$

- $\text{Left}(i) = 2i$

- $\text{Right}(i) = 2i+1$



- Nearly complete binary tree
 - Completely filled on all levels except possibly the last.
 - Last level filled from left up to a point.
 - Satisfies heap property
- Two kinds:
 - Min heap (largest element at the root)
 - Max heap (smallest element at the root)




The heap property

- Max-heap property: For every node i in the tree other than the root, $A[\text{parent}(i)] \geq A[i]$
- Min-heap property: ?

Max heap is used for heap sort

Height

- Height of a node in a heap: number of edges on the longest path downward to the leaf
- Height of the heap: height of the root



Let's do it

- What is the minimum and maximum number of elements in a heap of height h ?
- Prove that the height of a heap is $O(\log n)$.
- If all elements in a max-heap are distinct, where will the smallest value reside?
- Is a sorted array (ascending) a min-heap?
- Is the array $\{23, 17, 14, 6, 13, 10, 1, 5, 7, 12\}$ a max-heap?

- Can you identify the indices of the leaf nodes in a heap? Note, it is given there are n elements in the heap.

- Can you identify the indices of the leaf nodes in a heap? Note, it is given there are n elements in the heap.

Ans: $\{\lfloor n/2 \rfloor + 1, \dots, n\}$

Proof:

- Show that all nodes with indices $\{\lfloor n/2 \rfloor + 1, \dots, n\}$ are childless \rightarrow they are leaves.
- Part 1: All nodes in the range of indices $\{\lfloor n/2 \rfloor + 1, \dots, n\}$ are not having children in this range.
- Part 2: Any node not having children does fall in this range $\{\lfloor n/2 \rfloor + 1, \dots, n\}$

- Part 1: All nodes in the range of indices $\{\lfloor n/2 \rfloor + 1, \dots, n\}$ are not having children in this range.
- Let i be a node in this range. It's children are present at index $2i$ and $2i+1$. Assuming, $i = \lfloor n/2 \rfloor + 1$, we have $2i = 2\lfloor n/2 \rfloor + 2 > n$. Because the left child of the smallest i in the range $\lfloor n/2 \rfloor + 1, \dots, n$ is $>n \Rightarrow$ all nodes in range $\lfloor n/2 \rfloor + 1, \dots, n$ are childless.
- Part 2: Any node not having children does fall in this range $\{\lfloor n/2 \rfloor + 1, \dots, n\}$
- Let i be a node with no children. That is, $2i$ and $2i+1 > n$. $\Rightarrow i > \lfloor n/2 \rfloor$. $\Rightarrow i \in \{\lfloor n/2 \rfloor + 1, \dots, n\}$


Maintaining the heap property


Maintaining the heap property

- A procedure is called : MAX-HEAPIFY
- Input to MAX-HEAPIFY - Array A and index i
- Function assumes left(i) and right(i) are roots of max-heaps, but A[i] might not be. Hence, it floats the value down until the tree rooted at index i becomes a max-heap.


MAX-HEAPIFY(A, i)

```
1   $l = \text{LEFT}(i)$ 
2   $r = \text{RIGHT}(i)$ 
3  if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$ 
4       $largest = l$ 
5  else  $largest = i$ 
6  if  $r \leq A.\text{heap-size}$  and  $A[r] > A[largest]$ 
7       $largest = r$ 
8  if  $largest \neq i$ 
9      exchange  $A[i]$  with  $A[largest]$ 
10     MAX-HEAPIFY( $A, largest$ )
```






(a)



(b)



(c)

Running time: MAX-HEAPIFY

- At a given subtree of size n rooted at node i , MAX-HEAPIFY needs $\theta(1)$ to fix the heap property in $A[i]$, $A[\text{left}(i)]$ and $A[\text{right}(i)]$ + recursive calls to MAX-HEAPIFY on a subtree rooted at one of the children of i .
- The children's subtrees each have a size at most $2n/3$ - the worst case occurs when the bottom level is exactly half full (why?)
- Therefore, recurrence relation for MAX-HEAPIFY: $T(n) \leq T(2n/3) + O(1)$
- Solving using master's theorem, the solution is: $O(\log n)$

Let's do it

- Run MAX-HEAPIFY(A,3) on $A = \{27, 17, 3, 16, 13, 10, 1, 5, 7, 12, 4, 8, 9, 0\}$
- Suppose A.heapSize is the number of elements in the heap. What will happen when you call MAX-HEAPIFY(A, i) for $i > A.\text{heapSize}/2$?
- Modify the heapify procedure for a min-heap/Write procedure for MIN-HEAPIFY.

Building the heap

Building the heap


- Two ways:
 - William's method
 - Floyd's method (used in book)

Building a heap - William's method

- Consider an input array: A
= { 23, 17, 5, 10, 11, 26}

```
Williams Algorithm: top down
while not end of array,
    if heap is empty,
        place item at root;
    else,
        place item at bottom of heap;
        while (child > parent)
            swap(parent, child);
        go to next array element;
end
```

- Consider an input array: $A = \{23, 17, 5, 10, 11, 26\}$



Running time: William's method

- The number of operations required depends only on the number of levels the new element must rise to satisfy the heap property, thus the insertion operation has a worst-case time complexity of $O(\log n)$
- If we do this for every node, then we have $O(n \cdot \log n)$ as the running time for William's method -> not very efficient.
- Efficiency approach proposed by Floyd.

Building a heap - Floyd's method

- The elements in the sub array $A[(\lfloor n/2 \rfloor + 1) \dots n]$ are all leaves and so each is a 1-element heap to begin with.
- This method runs the MAX-HEAPIFY procedure on all the remaining nodes.
- It is a bottom-up approach.

BUILD-MAX-HEAP(A)


1 $A.\text{heapSize} = A.\text{length}$

2 for $i = \lfloor A.\text{length}/2 \rfloor$ down to 1


3. MAX-HEAPIFY(A, i)

A


4	1	3	2	16	9	10	14	8	7
---	---	---	---	----	---	----	----	---	---




(a)





(b)



(c)



(d)



Correctness of Heap sort

- Depends on the correctness of BUILD-MAX-HEAP procedure.

Correctness of BUILD-MAX-HEAP

- LOOP INVARIANT ??

Correctness of BUILD-MAX-HEAP

- LOOP INVARIANT
- **At the start of each iteration of the for loop, each node $i+1, i+2, \dots, n$ is the root of a max-heap.**
- Need to show that this invariant is true prior to the first loop iteration, that each iteration of the loop maintains the invariant, and the invariant provides a useful property to show the correctness when the loop terminates.

Initialisation:

- Prior to the first iteration of the loop $i = \lfloor n/2 \rfloor$. Each node $\lfloor n/2 \rfloor + 1, \lfloor n/2 \rfloor + 2, \dots, n$ is a leaf and is thus the root of a trivial max heap

BUILD-MAX-HEAP(A)

```
1 A.heapSize = A.length
2 for i =  $\lfloor A.length/2 \rfloor$  down to 1
3.  MAX-HEAPIFY(A, i)
```

Maintenance:

- The children of node i are numbered higher than i .
- By the loop invariant, they are both roots of max-heaps.
- When we call $\text{MAX-HEAPIFY}(A, i)$, it will make i a max-heap root. Moreover, it preserves the property that $i+1, i+2, \dots, n$ are all roots of max-heaps.
- Decrementing i in the floor loop update reestablishes the loop invariant for the next iteration.

BUILD-MAX-HEAP(A)

```
1  $A.\text{heapSize} = A.\text{length}$   
2 for  $i = \lfloor A.\text{length}/2 \rfloor$  down to 1  
3.  $\text{MAX-HEAPIFY}(A, i)$ 
```

BUILD-MAX-HEAP(A)

1 A.heapSize = A.length

2 for i = $\lfloor A.length/2 \rfloor$ down to 1

3. MAX-HEAPIFY(A, i)

Termination:

- At termination, $i=0$
- By the loop invariant each node 1, 2, ..., n is the root of a max-heap. Node 1 is.


Running time: Floyd's method

Simpler to understand:

- Each call to MAX-HEAPIFY requires $O(\log n)$ time.
- We make $O(n)$ such calls.
- Total = $O(n \cdot \log n)$.
- Although correct, but not a tight bound.

Running time: Floyd's method

- The time for MAX-HEAPIFY is not always $\log n$.
- It varies with the height of the node in the tree, and height of most nodes is small.
- The height of a n -node heap has a height of $\lfloor \log n \rfloor$ and at most $\lceil n/2^{h+1} \rceil$ nodes of any height h .



- The time required by MAX-HEAPIFY when called on a node of height h is $O(h)$.
- Thus, we can express the total time required by BUILD-MAX-HEAP as below:

- $$\sum_{h=0}^{\lfloor \log n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil O(h)$$

- $$O\left(n \sum_{h=0}^{\lfloor \log n \rfloor} \frac{h}{2^h}\right)$$

- We are already familiar with the series
$$\sum_{h=0}^{\infty} \frac{h}{2^h} = \frac{1/2}{(1 - 1/2)^2} = 2$$

- Thus we have,
$$O\left(n \sum_{h=0}^{\lfloor \log n \rfloor} \frac{h}{2^h}\right) = O\left(n \sum_{h=0}^{\infty} \frac{h}{2^h}\right) = O(n)$$

Heap can be built in linear time



Let's do it

- Build a heap from the array $A = \{5, 3, 17, 10, 84, 19, 6, 22, 9\}$

Pulling it all together - The heap sort algorithm

Heap Sort Algorithm

HeapSort(A)

1 BUILD-MAX-HEAP(A)

2 for i = A.length down to 2

3 Exchange A[1] with A[i]

4. A.heapSize = A.heapSize-1

5. MAX-HEAPIFY(A, 1)

HeapSort(A)


1 BUILD-MAX-HEAP(A)


2 for i = A.length down to 2

3 Exchange A[1] with A[i]

4. A.heapSize = A.heapSize-1


5. MAX-HEAPIFY(A, 1)






(b)

$$\hat{i} = 10$$




(c)

$i = 9$




$$i = 8$$




(e)


$i = 7$

$$i = 6$$



$i = 5$







(h)

$$i = 4$$

$$i = 3$$



$$i = 2$$



A

1	2	3	4	7	8	9	10	14	16
---	---	---	---	---	---	---	----	----	----

(1)

Priority Queues

Priority Queues

- Priority queue is a DS for maintaining a set S of elements each with an associated value called key.
- Types:
 - Min priority queue (implemented using min heaps)
 - Max priority queue (implemented using max-heaps)

Max priority queue

- Supports the following functions:
 - INSERT(S, x): Inserts an element x into set S (equivalent to the operation $S \cup \{x\}$)
 - MAXIMUM(S): returns element of S with largest key
 - EXTRACT-MAX(S): removes and returns the element of S with the largest key
 - INCREASE-KEY(S, x, k): increases the value of x's key to the new value k (assumed to be as large as x's current key value)

HEAP-MAXIMUM(A)

HEAP-MAXIMUM(A)

1 **return** $A[1]$

Time complexity?

HEAP-MAXIMUM(A)

HEAP-MAXIMUM(A)

1 **return** $A[1]$

Time complexity? $\Theta(1)$

HEAP-EXTRACT-MAX(A)

HEAP-EXTRACT-MAX(A)

```
1  if  $A.heap-size < 1$ 
2      error “heap underflow”
3   $max = A[1]$ 
4   $A[1] = A[A.heap-size]$ 
5   $A.heap-size = A.heap-size - 1$ 
6  MAX-HEAPIFY( $A, 1$ )
7  return  $max$ 
```

Time complexity?

HEAP-EXTRACT-MAX(A)

HEAP-EXTRACT-MAX(A)

```
1  if  $A.heap-size < 1$ 
2      error “heap underflow”
3   $max = A[1]$ 
4   $A[1] = A[A.heap-size]$ 
5   $A.heap-size = A.heap-size - 1$ 
6  MAX-HEAPIFY( $A, 1$ )
7  return  $max$ 
```

Time complexity?
 $O(\log n)$

HEAP-INCREASE-KEY(A, i, key)

HEAP-INCREASE-KEY(A, i, key)

```
1  if  $\text{key} < A[i]$ 
2      error “new key is smaller than current key”
3   $A[i] = \text{key}$ 
4  while  $i > 1$  and  $A[\text{PARENT}(i)] < A[i]$ 
5      exchange  $A[i]$  with  $A[\text{PARENT}(i)]$ 
6       $i = \text{PARENT}(i)$ 
```


Time complexity?

HEAP-INCREASE-KEY(A, i, key)


HEAP-INCREASE-KEY(A, i, key)

```
1  if  $\text{key} < A[i]$ 
2      error "new key is smaller than current key"
3   $A[i] = \text{key}$ 
4  while  $i > 1$  and  $A[\text{PARENT}(i)] < A[i]$ 
5      exchange  $A[i]$  with  $A[\text{PARENT}(i)]$ 
6       $i = \text{PARENT}(i)$ 
```


Time complexity?
 $\rightarrow O(\log n)$




(a)



(b)



(c)



(d)

MAX-HEAP-INSERT(*A*, *key*)

MAX-HEAP-INSERT(*A*, *key*)

- 1 $A.heap-size = A.heap-size + 1$
- 2 $A[A.heap-size] = -\infty$
- 3 HEAP-INCREASE-KEY(*A*, *A.heap-size*, *key*)

Time complexity?

MAX-HEAP-INSERT(*A*, *key*)

MAX-HEAP-INSERT(*A*, *key*)

1 $A.heap-size = A.heap-size + 1$

2 $A[A.heap-size] = -\infty$

3 HEAP-INCREASE-KEY(*A*, *A.heap-size*, *key*)

Time complexity?

$O(\log n)$